

RUSHIKESH S. WARE

rushikeshware@gmail.com | rware3@gmu.edu | (571) 435-1707 | <https://www.linkedin.com/in/rushikesh-ware10>

EDUCATION

George Mason University,

Master's in Data Analytics Engineering

Coursework: Big Data, Machine Learning, NLP with deep learning, Predictive Analytics

Expected Dec 2025

GPA: 3.93

Savitribai Phule Pune University, Pune, India

Bachelor of Engineering in Mechanical Engineering

May 2021

GPA: 3.52

WORK EXPERIENCE

Data Analyst

Cognizant Technology Solutions

Pune, India

Feb 2022 – Nov 2023

- Designed and optimized BigQuery SQL workflows that improved data delivery time by 25% and reduced compute cost by 18%.
- Refactored 10+ Python-based ETL pipelines to enhance model input accuracy by 12%, increasing downstream forecast precision.
- Automated CI/CD for data workflows using Docker, Kubernetes, and Airflow, reducing deployment incidents by 20% and saving 30+ hours/sprint.
- Collaborated with 6 cross-functional product teams to align analytics dashboards with user KPIs, increasing stakeholder adoption by 35%.
- Introduced documentation and code review standards across 4 teams, improving delivery consistency and reducing onboarding time by 30%.

Data Engineer

Forest Analytics (MCP x ClimateGPT)

Fairfax, VA

Aug 2025 – Dec 2025

- Architected and deployed a natural language SQL query pipeline over 500K+ environmental records, enabling zero-code querying by non-technical users.
- Built and maintained 6+ scalable data models for forest loss, CO₂ emissions, and land cover, reducing query latency by 70%.
- Integrated with ClimateGPT via API and Claude Desktop, generating automated explanations for 100+ test cases with 90% accuracy.
- Owned data quality validation across 4 metadata dimensions (schema match, value ranges, unit consistency, null handling), improving pipeline reliability by 85%.
- Containerized and deployed the system with Docker, cutting deployment time by 40% compared to manual configuration.

Data Analyst

NLP & SQL Transformer | [GitHub](#)

Fairfax, VA

Aug 2024 – Apr 2025

- Engineered a SQL generation engine using BART and schema-aware tokenization, achieving 45.6% exact match on Spider benchmark across 5 domains.
- Optimized inference runtime by 60% using ONNX acceleration and FastAPI, maintaining <200 ms response time across 10K test queries.
- Containerized the model pipeline with Docker and orchestrated with Kubernetes across 3 environments (dev, test, prod).
- Instrumented performance dashboards with Prometheus + Grafana, identifying 3 bottlenecks that reduced memory usage by 28%.

Data Analyst

Public Health Analytics | [GitHub](#)

Fairfax, VA

Jan 2024 – May 2025

- Analyzed** 20 years of overdose data using AWS EC2, Glue DataBrew, and Python regression techniques to uncover mortality trends.
- Highlighted** demographic risk factors, establishing fentanyl's disproportionate impact on ages 25–34.
- Configured** interactive dashboards and heatmaps to support evidence-based public health initiatives.

Data Analyst

Infrastructure Analytics & Optimization |

Fairfax, VA

Aug 2024 – Dec 2024

- Forecasted** structural ratings for 7,500+ bridges with XGBoost and Random Forest, yielding 45% R² accuracy.
- Incorporated** environmental and financial indicators to flag high-risk assets and guide maintenance decisions.
- Applied** Chi-Square, ANOVA, and NPV modeling to inform cost-effective regional investment strategies.

Data Analyst

FIFA 22 Data Analysis Project

Fairfax, VA

Jan 2024 – May 2025

- Conducted an in-depth analysis of FIFA 22 player performance trends using Python and R, uncovering key insights and achieving a 15% improvement in model accuracy (see report).
- Generated interactive visualizations and a comprehensive report that boosted data accessibility and stakeholder engagement by 20%.
- Transformed raw data into actionable information using advanced analytical tools, resulting in a 12% increase in performance insight efficacy.

Data Analyst

Cybersecurity Risk Modeling (KEV)

Fairfax, VA

Jan 2025 – May 2025

- Developed** an XGBoost model ingesting CVE, CVSS, EPSS, and CPE data; addressed class imbalance using SMOTE, PCA, and threshold calibration.
- Configured** cost-sensitive loss functions to improve model recall, achieving AUC = 0.996 and full recall for exploited vulnerabilities.
- Delivered** a deployable triage tool that integrates into security workflows for automated patch prioritization.

SKILLS AND TOOLS

Languages & Tools: Python, R, SQL (PostgreSQL, MySQL), Bash, Spark, Hadoop, Docker, Kubernetes, Airflow, Terraform, Git, ONNX, MLflow, Streamlit, Power BI, Tableau, Data Modeling, BigQuery, Kubernetes, Excel, MATLAB, PyTorch, FastAPI, Flask, TensorRT, TFLite, containerization, REST APIs, model quantization & pruning, automated testing, protobuf, Avro.

Cloud & Platforms: GCP, AWS (S3, EC2, Glue), Databricks, Snowflake, Looker, serverless deployments.

Techniques: EDA, Feature Engineering, Regression, Classification, Time Series, NLP, Deep Learning, SMOTE, PCA, Data Visualization, network architecture optimization, model benchmarking, observability (Prometheus, Grafana).